

```
In [ ]: import cv2
import numpy as np
from google.colab.patches import cv2_imshow
```

```
In [ ]: from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```
In [ ]: %cd /content/gdrive/My Drive/ENPM673/Project 2
/content/gdrive/My Drive/ENPM673/Project 2
```

## Task 1.

Design a video processing pipeline to find the 4 corners of the paper using Edge Detection, Hough Transform and Corner Detection.

Overlay the boundary(edges) of the paper and highlight the four corners of the paper. The boundary(edges) of the paper and highlight the four corners of the paper. (Note: that you must remove any frames which are blurry by using Variance of Laplacian and report the number of blurry frames removed). Example Pipeline:

1. Read the video and extract individual frames using OpenCV.
2. Skip blurry frames (use Variance of the Laplacian and decide a suitable threshold) Note: Any value below 150 for the Variance of the Laplacian, suggests that it's a blurry image.
3. Segment out the unwanted background area (example: convert to gray scale and keep white regions)
4. Detect edges pixels in each frame (you can use any edge detector)
5. Use the detected edge pixels to extract straight lines in each frame (hint: use Hough Transform)
6. Filter out “short” lines and only keep a few dominant lines.
7. Compute the intersections of the Hough Lines – these are the putative corners of the paper.
8. Verify the existence of those corners with Harris corner detector. (use OpenCV built-in function)
9. Filter out remaining extraneous corners that are not the 4 corners of the paper.
10. Generate the output video in which you have to overlay the 4 blue boundary lines and 4 red corners of the paper in each frame (excluding the blurry frames).

```
In [91]: # Read the video and extract individual frames using OpenCV

path = './videos/proj2_v2.mp4'
cap = cv2.VideoCapture(path)
if not cap.isOpened():
    print("ERROR: Could not open the video")
    exit()

# Laplacian function for Step 2
def laplacian(gray):
    laplacian_kernel = np.array([[0, 1, 0],
                                [1, -4, 1],
                                [0, 1, 0]], dtype=np.float64)
    laplacian = cv2.filter2D(gray, -1, laplacian_kernel)
    return laplacian

# Intersection function for Step 7
def find_intersection(line1, line2):
    x1, y1, x2, y2 = line1
    x3, y3, x4, y4 = line2

    denom = (x1 - x2) * (y3 - y4) - (y1 - y2) * (x3 - x4)
    if denom == 0:
        return None # parallel lines
    px = ((x1*y2 - y1*x2) * (x3 - x4) - (x1 - x2) * (x3*y4 - y3*x4)) /
    py = ((x1*y2 - y1*x2) * (y3 - y4) - (y1 - y2) * (x3*y4 - y3*x4)) /
    return (px, py)

# Initialize Video Writer
fourcc = cv2.VideoWriter_fourcc(*'mp4v') # Define the codec
output_path = './videos/output_proj2_v2.mp4'
frame_width = int(cap.get(3))
frame_height = int(cap.get(4))
out = cv2.VideoWriter(output_path, fourcc, 20.0, (frame_width, frame_height))
c = 0
while True:
    ret, frame = cap.read()
    if not ret:
        print("Could not receive the frame.")
        break

    #----- Step 2 -----
    # Skip blurry frames (use Variance of the Laplacian and decide a suitable threshold)
    # Note: Any value below 150 for the Variance of the Laplacian, suggests a blurry frame
    gray_img = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    var = laplacian(gray_img).var()

    if var<100:
        print("it's a blurry image, skipping...")
        c+=1
        continue

    #----- Step 3 -----
    # Detect lines in the frame
    edges = cv2.Canny(frame, 50, 150)
    lines = cv2.HoughLinesP(edges, 1, np.pi/180, 100, minLineLength=50, maxLineGap=10)
    if lines is not None:
        for line in lines:
            x1, y1, x2, y2 = line[0]
            cv2.line(frame, (x1, y1), (x2, y2), (0, 255, 0), 2)

    #----- Step 4 -----
    # Find intersection points of the lines
    intersections = []
    for i in range(len(lines)):
        for j in range(i+1, len(lines)):
            line1 = lines[i][0]
            line2 = lines[j][0]
            px, py = find_intersection(line1, line2)
            if px is not None:
                intersections.append((px, py))

    #----- Step 5 -----
    # Draw intersection points
    for intersection in intersections:
        cv2.circle(frame, (int(intersection[0]), int(intersection[1])), 5, (255, 0, 0), -1)

    #----- Step 6 -----
    # Write the frame to the output video
    out.write(frame)

    #----- Step 7 -----
    # Show the frame
    cv2.imshow('Frame', frame)

    #----- Step 8 -----
    # Check for user input
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

    #----- Step 9 -----
    # Print progress
    print(f'{c} frames processed')

#----- Step 10 -----
# Release resources
cap.release()
out.release()
cv2.destroyAllWindows()
```

```
else:

#----- Step 3 -----
# Segment out the unwanted background area (example: convert to gray s
    # x, thresh = cv2.threshold(gray_img,220,255, cv2.THRESH_BINARY)
    y = cv2.GaussianBlur(gray_img,(5,5), 0)
    thresh = np.where(y>220,255,0).astype(np.uint8)

#----- Step 4 -----
# Detect edges pixels in each frame (you can use any edge detector)
edge = cv2.Canny(thresh,100,200)

#----- Step 5 & 6 -----,
# Use the detected edge pixels to extract straight lines in each frame
# Filter out "short" lines and only keep a few dominant lines
lines = cv2.HoughLinesP(edge, 1, np.pi / 180, threshold=60, minLineLength=50, maxLineGap=10)
for line in lines:
    x1, y1, x2, y2 = line[0]
    cv2.line(frame, (x1, y1), (x2, y2), (255, 0, 0), 2)

#----- Step 7 -----
# Compute the intersections of the Hough Lines – these are the putative
# intersection points
intersections = []
for i in range(len(lines)):
    for j in range(i+1, len(lines)):
        intersection = find_intersection(lines[i][0], lines[j][0])
        if intersection: # Ensure the intersection point is valid
            intersections.append(intersection)
            # For visualization, mark the intersection
            # cv2.circle(frame, (int(intersection[0])), int(intersection[1]), 2)

#----- Step 8 -----
# Verify the existence of those corners with Harris corner detector. (This
# step is optional, but it helps to filter out false positives)
#----- Step 9 -----
# Filter out remaining extraneous corners that are not the 4 corners of the
# rectangle
dest = cv2.cornerHarris(gray_img, 2, 3, 0.04)
corners = np.argwhere(dest>0.5*dest.max())
# dest_dilated = cv2.dilate(dest, None) # Results are marked through dilation
# frame[dest_dilated > 0.01 * dest_dilated.max()]=[255, 0, 0] # Red color
# Verify intersections with Harris corners
verified_intersections = []
for intersection in intersections:
    for corner in corners:
        if np.linalg.norm(intersection - corner[::-1])<20:
            verified_intersections.append(intersection)
for i in verified_intersections:
    cv2.circle(frame, (int(i[0]),int(i[1])), 5, (0, 0, 255), -1)

#----- Step 10 -----
```

```
# Generate the output video in which you have to overlay the 4 blue bc  
# # and 4 red corners of the paper in each frame (excluding the blurry  
out.write(frame)  
  
cv2.imshow(frame)  
print(c, ' frames skipped. ')  
if cv2.waitKey(1) == ord('q'):  
    break  
out.release() # Finalize the video file  
cap.release()  
cv2.destroyAllWindows()
```



## Task 2.

## Part A

Given four overlapping images of a far-way building taken from the same camera (images may be taken with both rotation and slight translation). Design an image processing pipeline to stitch these images to create a panoramic image. (Note: You can consider far away objects/features to be approximately on the same plane.)

## Part B

In general, why does panoramic mosaicing work better when the camera is only allowed to rotate at its camera center?

Example pipeline:

1. Extract features from each frame (You can use any feature extractor and justify).
2. Match the features between each consecutive image and visualize them. (hint: Use RANSAC)
3. Compute the homographies between the pairs of images.
4. Combine these frames together using the computed homographies.

```
In [ ]: #import libraries
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
import random
from itertools import groupby, product
from scipy.spatial.transform import Rotation as Ro
from google.colab.patches import cv2_imshow

# Function for extracting the features
def feature_extractor(img1, img2):
    sift = cv.SIFT_create() # Create a SIFT object
    kp1, ds1 = sift.detectAndCompute(img1, None)
    kp2, ds2 = sift.detectAndCompute(img2, None)
    return kp1, ds1, kp2, ds2 # Return keypoints and descriptors for both images

# Function for matching the features
def feature_matcher_1_2(k1, k2, d1, d2, img1, img2):
    FLANN_INDEX_KDTREE = 1
    index_params = dict(algorithm = FLANN_INDEX_KDTREE, trees = 5)
    search_params = dict(checks=50)
    flann = cv.FlannBasedMatcher(index_params,search_params) # Create FLANN matcher
    matches = flann.knnMatch(d1,d2, k =2)
```

```
good_matches = []# List to store good matches

for m, n in matches:
    if m.distance < 0.8 * n.distance:
        good_matches.append(m)

top_matches = matches[:50]

img12 = cv.drawMatchesKnn(img1,k1,img2,k2,top_matches,None)
src_pts = np.float32([k1[m.queryIdx].pt for m in good_matches]).reshape(-1,1,2)
dst_pts = np.float32([k2[m.trainIdx].pt for m in good_matches]).reshape(-1,1,2)
klist= []
for mat in good_matches:

    (x1, y1) = k1[mat.queryIdx].pt
    (x2, y2) = k2[mat.trainIdx].pt

    klist.append([(x1, y1), (x2, y2)])# Append point pairs to the list
plt.imshow(img12)
plt.show()
return top_matches, img12, klist, src_pts, dst_pts

# Function for cropping the black part in the stitched image
def trim_black_part(img):

    if not np.sum(img[0]):
        return trim_black_part(img[1:])
    elif not np.sum(img[-1]):
        return trim_black_part(img[:-1])
    elif not np.sum(img[:,0]):
        return trim_black_part(img[:,1:])
    elif not np.sum(img[:, -1]):
        return trim_black_part(img[:, :-1])

    return img

image_1 = cv.imread("./videos/img1.jpg", cv.IMREAD_COLOR)
image_2 = cv.imread("./videos/img2.jpg", cv.IMREAD_COLOR)
image_3 = cv.imread("./videos/img3.jpg", cv.IMREAD_COLOR)
image_4 = cv.imread("./videos/img4.jpg", cv.IMREAD_COLOR)
# image 3 and 4
# 1) Extract features from each frame (You can use any feature extractor)
key_3, des3, key_4, des4 = feature_extractor(image_3, image_4)

# 2) Match the features between each consecutive image and visualize them
top, image_34, key_list_34, src_pts, dst_pts = feature_matcher_1_2(key_3, des3, key_4, des4)

# 3) Compute the homographies between the pairs of images.
H_34, mask = cv.findHomography(src_pts, dst_pts, cv.RANSAC)
```

```
# 4) Combine these frames together using the computed homographies.
H_43 = np.linalg.inv(H_34)
width_3 = image_3.shape[1]
height_3 = image_3.shape[0]
width_4 = image_4.shape[1]

result_34 = cv.warpPerspective(image_4, H_43, (width_3 + width_4, height_3))
result_34[0:image_3.shape[0], 0:image_3.shape[1]] = image_3

result_34 = trim_black_part(result_34)

cv2.imshow(result_34)
#image 2 and 3
# 1) Extract features from each frame (You can use any feature extractor)
key_2, des2, key_3, des3 = feature_extractor(image_2, image_3)

# 2) Match the features between each consecutive image and visualize them
top, image_23, key_list_23, src_pts, dst_pts = feature_matcher_1_2(key_2, des2, key_3, des3)

# 3) Compute the homographies between the pairs of images.
H_23, mask = cv.findHomography(src_pts, dst_pts, cv.RANSAC)

# 4) Combine these frames together using the computed homographies.
H_32 = np.linalg.inv(H_23)
width_2 = image_2.shape[1]
height_2 = image_2.shape[0]
width_3 = result_34.shape[1]

result_23 = cv.warpPerspective(result_34, H_32, (width_2 + width_3, height_2))
result_23[0:image_2.shape[0], 0:image_2.shape[1]] = image_2

result_23 = trim_black_part(result_23)
cv2.imshow(result_23)

#image 1 and 2

# 1) Extract features from each frame (You can use any feature extractor)
key_1, des1, key_2, des2 = feature_extractor(image_1, image_2)

# 2) Match the features between each consecutive image and visualize them
top, image_12, key_list_12, src_pts, dst_pts = feature_matcher_1_2(key_1, des1, key_2, des2)

# 3) Compute the homographies between the pairs of images.
H_12, mask = cv.findHomography(src_pts, dst_pts, cv.RANSAC)

# 4) Combine these frames together using the computed homographies.
H_21 = np.linalg.inv(H_12)
width_1 = image_1.shape[1]
```

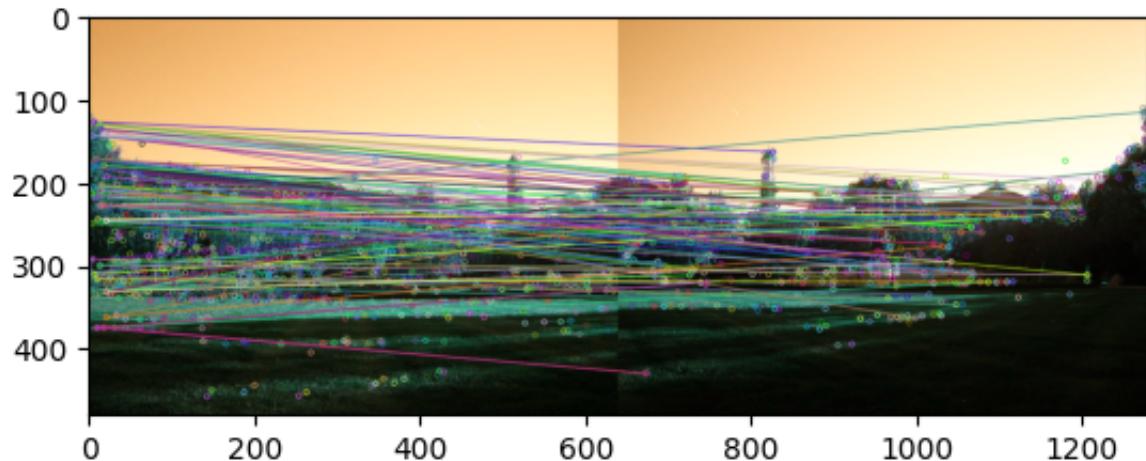
```
height_1 = image_1.shape[0]
width_2 = result_23.shape[1]

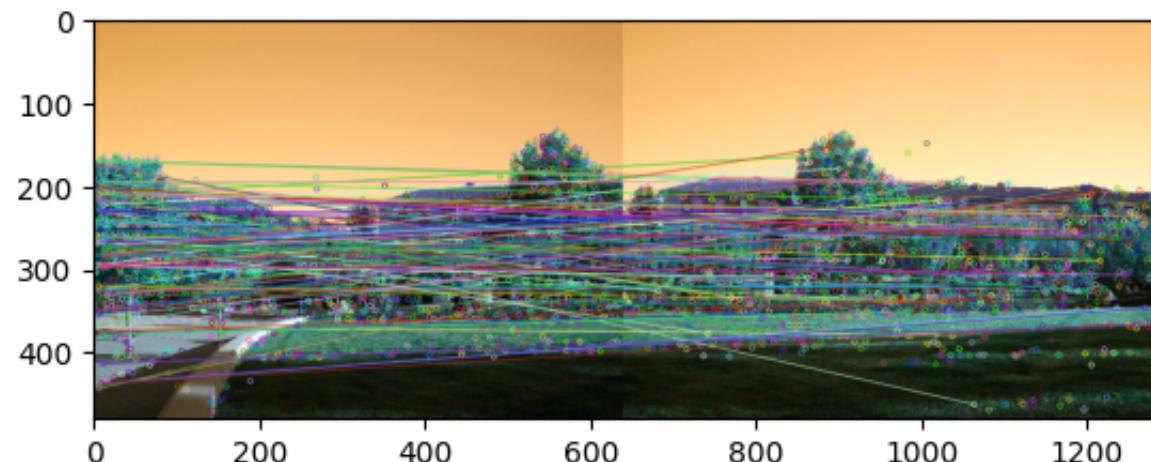
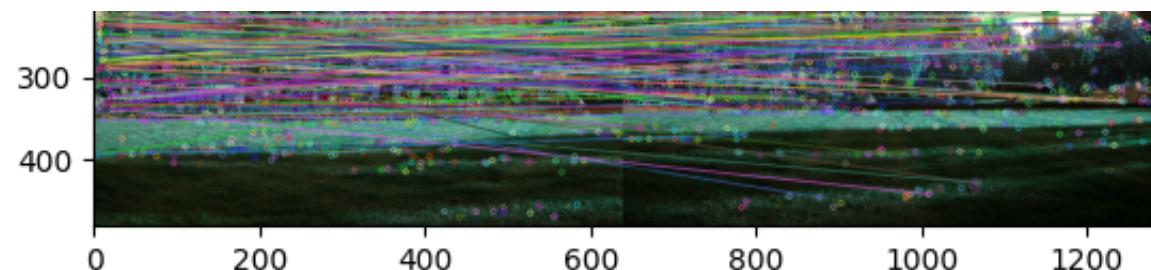
result_12 = cv.warpPerspective(result_23, H_21, (width_1 + width_2, height_1))

result_12[0:image_1.shape[0], 0:image_1.shape[1]] = image_1

result_12 = trim_black_part(result_12)
cv2_imshow(result_12)

cv.waitKey() & 0xFF == ord("q")
cv.destroyAllWindows()
```





In [ ]:

